



AI agents that work under your rules.

Nexus is an ERP for industrial SMEs. Its AI agents work under the same permissions, approvals and audit trail as your staff.

September 2026



Sales Operations Agent

L1 Recommend

3 quotes have had no reply for 7 days. I propose a follow-up task for each.

TKL-2026-000184 · 7 days

TKL-2026-000191 · 8 days

TKL-2026-000203 · 9 days

Is this suggestion right?

Right

Wrong

Problem

About three in four small Turkish firms work without an ERP.

Firms with 10-49 employees

23.6%

use ERP software in Türkiye.

EU average

41%

Firms with 50-249 employees

46.0%

use ERP software in Türkiye.

EU average

65%

Firms with 10+ employees

7.5%

use AI in Türkiye.

EU average

20.0%

Türkiye trails the EU on ERP by 17 points for small firms and 19 points for medium firms.
That gap sits in the firms Nexus is built for.

What it costs

An unanswered quote, a late collection, a forgotten follow-up: each one is lost revenue.

Unanswered quote

The quote went out 12 days ago and nobody replied. It is about to expire and no one noticed.

▲ Lost sale

Late collection

An overdue invoice reaches no one's desk until the month-end report.

▲ Cash tied up

Forgotten follow-up

The call-back promised in a meeting stays in a diary. The customer calls a competitor.

▲ Lost customer

About three in four Turkish firms with 10-49 employees work without an ERP, so these follow-ups depend on one person remembering.

Why now

Agents are reaching business software faster than the controls around them.

40%

of enterprise apps will include task-specific AI agents by the end of 2026, up from under 5% in 2025.

Gartner, 26 Aug 2025

40%

of enterprises will demote or decommission autonomous agents by 2027 because of governance gaps.

Gartner, 26 May 2026

Aug 2026

Logo, the local market leader, said it is deepening its agentic AI capabilities.

Logo H1 2026 results

Türkiye published its AI Action Plan 2026-2030 on 18 August 2026.

The principle

Three rules hold for every agent in Nexus.

Rule 1

An agent is a user.

It has its own role and sees only what that role allows. It cannot sign in, and it can never approve its own request.

Rule 2

The model never touches the database.

The agent works only through the same controlled actions a person uses. Anything that reaches a customer waits for a person.

Rule 3

Autonomy grows only with evidence.

The agent starts in shadow mode, where admins rate every suggestion. An admin raises its level on that evidence.

What the agent does

The agent checks six situations every morning.

1 Quote about to expire

Finds

Sent quotes whose validity is about to end.

At L2 and above

A follow-up task. At L3, a revision or a reminder e-mail.

2 Unanswered quote

Finds

Quotes sent 7 or more days ago with no reply.

At L2 and above

A follow-up task. At L3, a reminder e-mail.

3 Overdue follow-up

Finds

Promised follow-ups that were not done on time.

At L2 and above

A notification to the follow-up's owner.

4 Accepted, not ordered

Finds

Accepted quotes that have not become an order.

At L2 and above

A follow-up task, minutes after acceptance.

5 Stalled project

Finds

Open projects with no activity for a long time.

At L2 and above

A follow-up task for the project's owner.

6 Overdue receivable

Finds


Invoices 7 or more days past due.

At L2 and above


A follow-up task. At L3, a payment reminder e-mail.

The model is called only when plain code has found something to judge.

The autonomy ladder stops where a person approves.



Gartner's four autonomy levels, 26 May 2026



Every run takes the same six steps, all on record.

- 1 Trigger**
Every day at 07:30, and minutes after a quote is accepted.
- 2 Find**
Plain code finds the records. Nothing to find: no model call, no cost.
- 3 Mask**
Names become random codes, amounts are rounded, dates become day counts.
- 4 One model call**
Through Cloudflare AI Gateway, priced before it is made, never inside a database transaction.
- 5 Act by level**
A suggestion at L1, a follow-up task at L2, an approval request at L3.
- 6 Record**
Model, tokens, cost and the summary as sent. Every action is audited under the agent's name.

 **Sales Operations Agent** L1 Recommend

3 quotes have had no reply for 7 days. I propose a follow-up task for each.

Sent: codes for the quote and company, a rounded amount, days waiting


Never sent: names, contacts, e-mail, phone

Claude · Amazon Bedrock · AI Gateway

Is this suggestion right? Right Wrong

What the model sees

Your customer's name never reaches the model.



In a breach, all anyone could read is "[C-7KQ]'s quote: 120,000 TRY, unanswered for 9 days". The key to the codes lives in Nexus's memory for one run, then it is dropped.

Never sent

Company and contact names

E-mail

Phone

Tax number

Address

Notes

Follow-up titles

Record ids

Every column the agent reads is classified. An unclassified column fails the tests.

Nothing reaches a customer without a person.

- 1 The agent asks**
At L3 it raises an approval request under its own rule: quote revision, quote reminder or payment reminder.
- 2 A rule decides who approves**
An admin configures the rule. Without one, the request falls back to an L2 follow-up.
- 3 A person sees everything**
The record, the agent's reasoning and exactly what approving will do.
- 4 Approval runs it as them**
The e-mail is our fixed Turkish text, signed with the business unit. The model never writes it.
- 5 Nothing else can**
The agent can never approve its own request. The decision is written to the audit trail.

■ Needs your decision

Payment reminder e-mail: Karaduman Tekstil A.Ş.

Requested by	Invoice
 Sales Operations Agent	FTR-2026-000664
Open amount	Due
₺94,250.00	4 Sep 2026 · 25 days overdue

An agent cannot approve its own request. This decision belongs to a person.

Approve

Reject

One page to steer the agent, and one switch to stop it.

Kill switch

Stops the agent and disables its account. It is checked before every action.

Level

L0 to L3, chosen by an admin. Every change is audited.

Cost and usage

Today's cost against the daily cap, and the week's allowance.

Run ledger

Every run with its verdicts, its cost and the exact summary sent to the model.

What each tool sends

Read from the tool's own definition, so it cannot claim more than the code does.

 **Sales Operations Agent**

Agent on ☒

Level

L0 L1 L2 L3

Today

This week

21% of the daily cap

Run ledger

● Today 07:30

Finished: 3 suggestions, 1 dismissal

○ Yesterday 07:30

Nothing to judge, no model call

◆ Mon 14:02

Stopped: over budget, not retried

Ten modules share one customer record, one permission model and one audit trail.

CRM

Accounts, contacts and follow-ups

Sales

Quotes, orders and deliveries

Procurement

Requests, purchase orders and approvals

Inventory

Stock and movements

Finance

Visibility over your accounting system

HR

Staff records and recruitment

Tasks

Work items across all modules

Assets

Company assets and assignments

Fleet

Vehicles, tracking and fuel

Custody

Items issued to staff

One isolated instance and database per company. Multilanguage and multicurrency support.

Model route

One exit for every model call, and three routes to the model.

Cloudflare

Job in Nexus

AI Gateway: the single exit for every model call. It holds the cloud keys.

What the customer gets

No AWS key on the server, and no copy of requests or responses at the gateway.

AWS

Job in Nexus

The primary route: Claude on Amazon Bedrock.

What the customer gets

Bedrock does not store prompts or use them for training.
Invocation logging stays off.

Microsoft Azure

Job in Nexus

Azure AI Foundry as a second route. Later, sign-in with Microsoft accounts.

What the customer gets

A second route today. Automatic fallback is planned.

Agent → Cloudflare AI Gateway → Claude on Amazon Bedrock • Anthropic direct and Azure AI Foundry selectable

The security layer was built before the agents.

Identity

- scrypt password hashing
- Two-step sign-in (TOTP)
- Lockout after 5 failed tries
- Sessions revoked on password change

Permissions

- Role and business-unit scope on every query
- No one approves their own request
 - The agent is a user that cannot sign in
 - Names never reach a model

Data

- One database per company
- Seven database roles; the app cannot touch audit records
- Chained, locked audit trail
- KVKK retention and erasure flows

Continuity

- Encrypted backup every 6 hours
- The server cannot read old backups
- At most 6 h of data lost, back in 4 h
- Restore drills with evidence

Also live on the platform: encrypted sessions (iron-session), HTTPS only with HSTS, a Content Security Policy, a client IP header that cannot be spoofed, containers per company and a 148-scenario browser test suite.

What happens if something goes wrong?

A password is stolen

Sign-in fails without the second-step code. Guessing locks after 5 failed tries. Resetting the password closes every open session.

An employee rewrites history

They touch only what their role allows. Not even the app can change audit records; if an old entry changes, the chain breaks and it shows.

The server is lost

Encrypted backups sit in an account separate from production. The target: working again within 4 hours, with at most 6 hours of data lost.

An agent makes a mistake

It works only through controlled actions and cannot approve its own request. Anything that reaches a customer waits for a person. One switch stops it and closes its account.

The model provider is breached

All that leaks is codes, rounded amounts and day counts: "[C-7KQ]'s quote: 120,000 TRY, unanswered for 9 days". The key never leaves Nexus.

The spending limit is reached

Each call is priced before it is made. If the daily cap or the weekly allowance would be passed, the agent does not run and the run is recorded in red.

Traction

Nexus is in production at two companies, and its first agent is built.

Production

2 companies

• Live

22 Sep 2026

Shadow mode

• Live

25-28 Sep 2026

First agent built

Six signals, shadow-mode ratings, follow-ups, approval requests and a kill switch.

Next: the first live model run, then a pilot at a partner company on real data.

Market

Our first market is about 89,000 Turkish manufacturers with 10 to 249 people.

€2.3bn

TAM: all Turkish firms with 10-249 employees

354,700 firms at list price. All Turkish enterprise application spend today is about USD 420 million.

€577m

Core SAM: manufacturers with 10-249 employees

About 88,700 firms

€3.7m

SOM: yearly recurring revenue at 500 companies

Year 5 target. 80% small and 20% medium firms.

Pricing is per module, with three AI tiers on top.

Module Based Pricing


A subscription per company, priced by module count.
Indicative range: €390 to €1,500+ a month.

AI Functions Pricing

Three AI tiers on top of the modules.
Each tier grants a weekly AI allowance.


Weekly reset

Users see the allowance as a percentage.
It resets every Monday at 00:00, and spending never goes past it.



Illustration

Global vendors ship governed agents and local vendors ship assistants.



What others already do

Business Central, NetSuite and Salesforce treat agents as users with roles and human approval.

Where Nexus differs

The combination: governed agents in Turkish, one instance per company, a shadow-mode ladder, customer names that never reach the model and a disclosure on every suggestion.

The main threat

Logo has 230,000+ customers and 1,000+ partners and is moving into agentic AI. Speed matters.

The team: four electrical engineers and one commercial lead.

MM

Melih Manisalıgil

BSc Electrical Engineering, Istanbul Technical University
Project lead and system architect.

OT

Onur Topuz

Sinology graduate, business development
Lead for commercialisation and market research.

RA

Raşit Arslan

Electrical and electronics engineer.
Owner of Karela.
R&D manager, responsible for the execution infrastructure.

EŞ

Eda Şahin

Electrical and electronics engineer
R&D engineer for security validation and negative testing.

HZ

Hamdi Zengin

Electrical and electronics engineer
R&D engineer for measurement, telemetry and the test setup.

The roadmap takes the agent from the demo to real data.

Now

• Live and built

- Nexus in production at two companies
- Authorization engine in shadow mode
- Sales Operations Agent built: L1 to L3, masked data, kill switch

Next

○ In progress

- First live model run through Cloudflare AI Gateway
 - A pilot at a partner company, in shadow mode on real data
 - Tier allowances and AI pricing
- Nexus's sor: questions answered within the asker's permissions

Later

○ Planned

- Automatic fallback between model routes
 - Sign-in with Microsoft; Türkiye hosting for every install
- MCP access for ChatGPT and Claude under the same rules

The ask

From an incubator we ask for pilots, cloud credits and mentors.

Pilots

Introductions to industrial SMEs and corporate partners that will run the agent in shadow mode on their own data.

Cloud credits

AWS Activate, Microsoft for Startups and Cloudflare for Startups, to run Bedrock, Azure AI Foundry and AI Gateway.

Mentors

Enterprise sales to industrial firms, AI governance and investor readiness.

The AI CoE asks for AI at the core of a working product.

What the program asks

- A working product or prototype with a validated use case
- AI at the core of the solution
- Rolling review, applications open until December 2026


What Nexus brings

- An ERP in production at two companies
- A Sales Operations Agent built on the same permission and approval model
- Masked data: customer names never reach the model

What we want from it

- Corporate pilots through Plug and Play's network
- Mentoring on AI safety evaluation and enterprise sales
- Access to ITU faculty to publish our evaluation results
- Our project lead is an ITU Electrical Engineering graduate

One isolated instance per company and one gateway for every model call.



Engineering evidence behind the platform and the agent.

179

database migrations in
production, owned per
module

148

browser smoke cases
in production

2,442

unit tests in the AI line,
at the first agent phase

10

modules in one
codebase

Stack

Next.js 16, React 19, PostgreSQL 16, Docker and Tailwind 4.
One isolated instance per company.

Design system

Status colours follow the colour meanings of IEC 60073, with
a shape for every state. Violet marks agent work and nothing
else.

Claim status: the live platform.

Claim	Status
In production at two companies	Live
Ten modules on one customer record, permission model and audit trail	Live
Inventory module	Pilot
Finance: visibility over the accounting system	Live
One isolated instance and database per company	Live
Role and business-unit scope on every query	Live
No one can approve their own request	Live
Tamper-evident, hash-chained audit trail and two-factor login	Live
Controlled retention and erasure flows for personal data	Live
Turkish and English; TRY, EUR and USD	Live

Claim status: the agent and the model route.

Claim	Status
Permission engine, approval inbox and audit chain the agent works under; one isolated instance and database per company	Live
The agent is a user with its own role; it cannot sign in, change its role or approve its own request	Built
It reads through six classified tools and writes only through guarded services	Built
Masked data: random codes, rounded amounts, day counts; names never sent	Built
Four autonomy levels: L2 follow-ups for the owner; L3 approval requests run as the approver	Built
Event triggers: the agent reacts within minutes to business events such as an accepted quote	Built
Customer e-mails are signed by the business unit	Built
Run ledger with cost, budgets that fail closed, kill switch and a Disclosure page	Built
Claude on Bedrock, Anthropic or Azure AI Foundry behind Cloudflare AI Gateway	Ready
OpenAI and on-premise local inference adapter	Next
Pilot on a company's real data	Next

What leaves the server when the agent calls a model.

Data	What the model receives
Company, project, quotation, invoice, business unit	A random code, made fresh for every run
Amounts	Rounded to two significant figures, with the currency
Dates	Day counts: days waiting, days overdue, days to expiry
Stage and signal	As they are
Names, e-mails, phones, tax ids, addresses	Never
Notes, follow-up titles, record ids, HR data	Never

The key from code to name lives in memory for one run and is dropped. Each run stores the summary exactly as sent, codes only.

Every column a tool reads is classified; an unclassified column fails the tests. One model call per run, never inside a database transaction.

TLS on every hop. No request or response body is logged. Gateway and Bedrock logging stay off.

Claim status: security.

Claim	Status
Encrypted, signed session cookie (iron-session), 12-hour lifetime	Live
Passwords hashed with scrypt	Live
Lockout after 5 failed tries per user, 20 per IP address	Live
Sessions revoked on password change or access change	Live
CSP, frame embedding blocked, HSTS, automatic HTTPS	Live
Client IP header rewritten by the proxy	Live
Seven database roles; the app cannot touch audit records	Live
Encrypted off-site backups every 6 hours, RPO 6 h, RTO 4 h	In build
Admin can stop the agent with one switch that also disables its account	Built

What each vendor ships, as of September 2026.

Vendor	Agents shipped	Roles and human approval	Per-run audit or cost	Türkiye hosting or on-premise AI
Logo, Mikro, Nebim, ETA, Uyumsoft	Assistants only	Not documented	Not documented	On-premise cores; AI hosting not documented
Microsoft Business Central	Yes	Yes	Yes	Turkish tenants in the UAE; agents lack Turkish
SAP Business One	No; AI-generated SQL	People validate the SQL	Not documented	Core on-premise; AI in SAP cloud
Odoo 19	Yes	Tools run unconditionally	Not documented	AI calls OpenAI or Gemini
Oracle NetSuite	Yes	Yes	Explanations	No Türkiye region, no on-premise
Salesforce Agentforce	Yes	Yes	Yes	No Türkiye region, no on-premise
Nexus	Built	Built	Built	Instance per company; Türkiye hosting in progress

How we sized the market.

Layer	Definition	Firms	Per year
TAM	All Turkish firms with 10-249 employees, at list price	354,700	€2.3bn
Core SAM	Manufacturers with 10-249 employees (25% share, 2020)	88,700	€577m
Extended SAM	Manufacturing, construction and real estate, wholesale and retail	262,000	€1.7bn
SOM, year 5	500 paying companies, 80% small and 20% medium	500	€3.7m ARR

$TAM = 306,000 \times €4,680 + 48,700 \times €18,000$

$SOM = 500 \times (0.8 \times €4,680 + 0.2 \times €18,000)$

Our entry price is about ten times Logo's average revenue per customer (about USD 526). Nexus is priced for firms with complex operations. The 2020 sector split will be replaced with TÜİK counts by sector and size class.